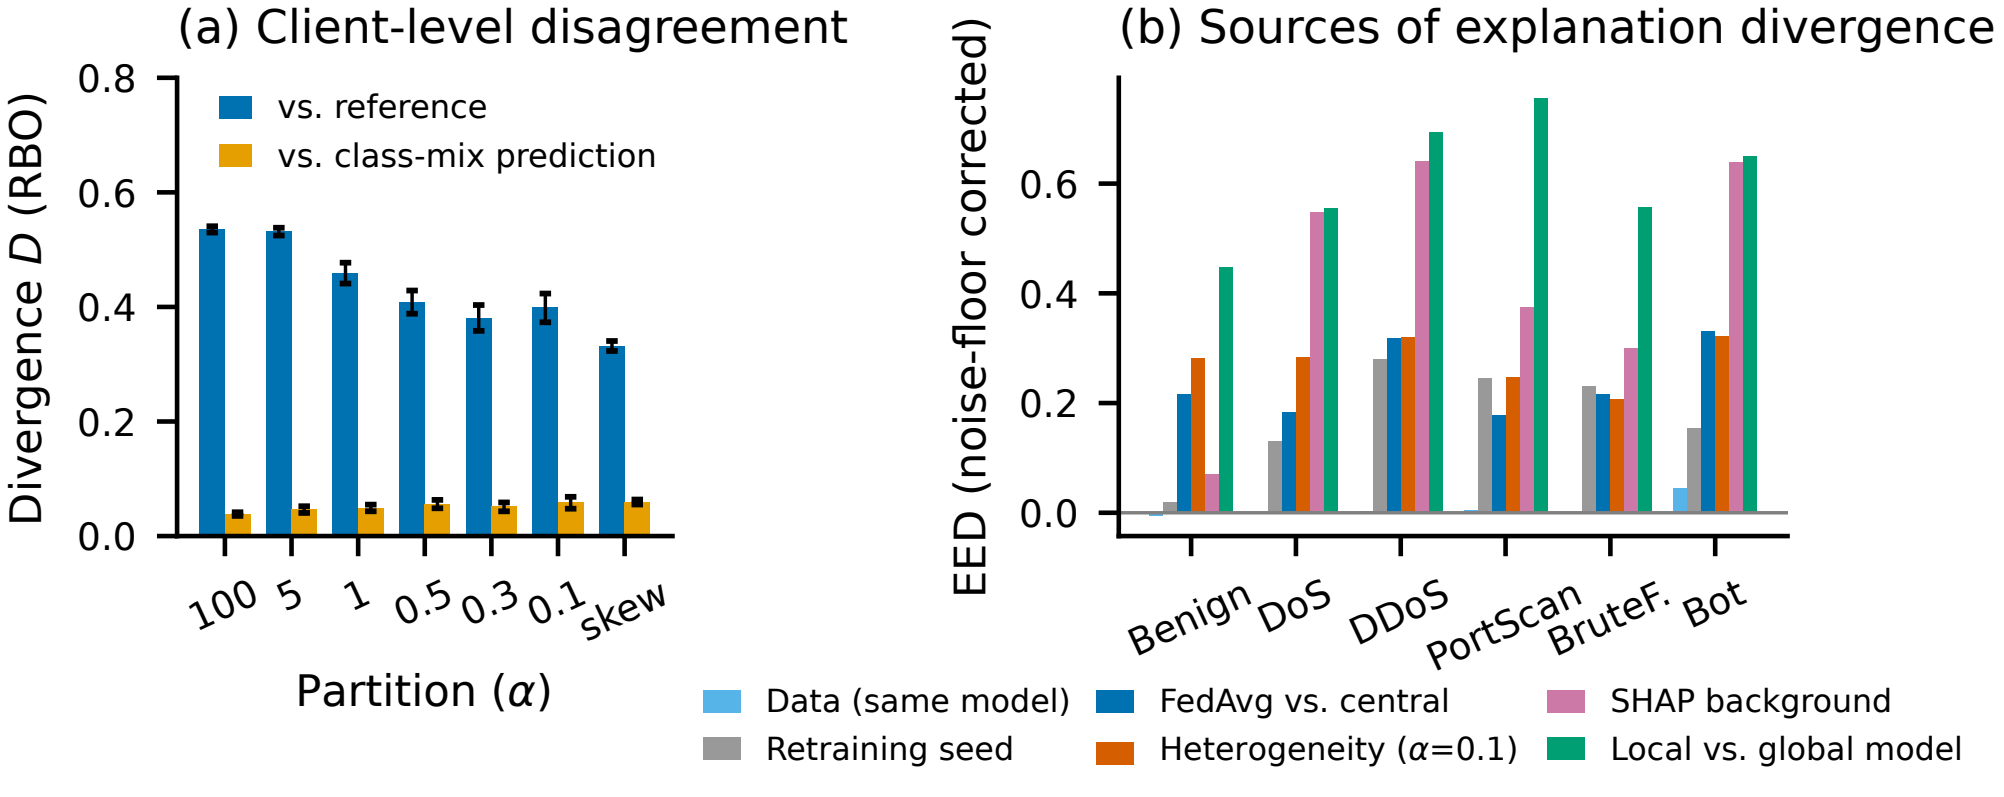
(a) Client-level disagreement
Divergence D (RBO)
0.8
0.6
0.4
0.2
0.0
vs. reference
vs. class-mix prediction
100
5
1
0.5
0.3
0.1
skew
Partition (α)
(b) Sources of explanation divergence
EED (noise-floor corrected)
0.6
0.4
0.2
0.0
Benign
DoS
DDoS
PortScan
BruteF.
Bot
Data (same model)
FedAvg vs. central
SHAP background
Retraining seed
Heterogeneity (α=0.1)
Local vs. global model